

IRE Assignment 1 – Text Indexing and Retrieval

Lakshay Baijal -- 2024202006

Github Link - https://github.com/LakshayBaijal/IRE_Assignments_Lakshay

1. Introduction

It focuses on building and analyzing a self implemented inverted index and comparing its query latency and performance with pre built indexing systems. The task involves preprocessing, index construction, query execution, and latency measurement on two datasets *Wikipedia* and *News*.

2. Preprocessing

The preprocessing step involves tokenizing, cleaning, and filtering raw text data before indexing.

```
• (venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_Hyderabad_CSIS/Semester_3/IRE
news.py
Processing News docs: 17it [00:00, 7604.86it/s]

✅ Done!
Processed file saved to: ../Dataset/webhose-news/data/processed/news_preprocessed.jsonl
Plots saved to:
→ plots/news_before.png
→ plots/news_after.png
```

Preprocessing results for News dataset

```
• (venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_Hyderabad_CSIS/Semester_3/IRE
wiki.py
Processing Wiki docs: 10000it [01:07, 147.65it/s]

✅ Done!
Processed file saved to: ../Dataset/Wiki_Dataset/data/processed/wiki_preprocessed.jsonl
Plots saved to:
→ plots/wiki_before.png
→ plots/wiki_after.png
```

Preprocessing results for Wiki dataset

4. Query Execution (Wiki Dataset)

Queries were executed against the Wiki index using the implemented search function.

```
○ (venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_Hyderabad_CSIS/Semester_7/Projects/Project_1/Python/te.py wiki_index.db
# then type queries (or 'exit' to quit)

Enter your query (or type 'exit' to quit): anarchism

🔍 Query Results for: anarchism
-----
1. Doc ID: 12 | Score: 1.000000 | length: 6843
2. Doc ID: 339 | Score: 1.000000 | length: 6383
3. Doc ID: 1023 | Score: 1.000000 | length: 8540
4. Doc ID: 2934 | Score: 1.000000 | length: 6115
5. Doc ID: 4035 | Score: 1.000000 | length: 9845
6. Doc ID: 4163 | Score: 1.000000 | length: 10182
7. Doc ID: 4874 | Score: 1.000000 | length: 9834
8. Doc ID: 5708 | Score: 1.000000 | length: 110
9. Doc ID: 5822 | Score: 1.000000 | length: 4305
10. Doc ID: 4049922 | Score: 1.000000 | length: 1509
-----

Enter your query (or type 'exit' to quit):
```

Output results for the anarchise Wiki query.

```
Enter your query (or type 'exit' to quit): computer science

🔍 Query Results for: computer science
-----
1. Doc ID: 742 | Score: 1.000000 | length: 138
2. Doc ID: 1134 | Score: 1.000000 | length: 2246
3. Doc ID: 1169 | Score: 1.000000 | length: 324
4. Doc ID: 1175 | Score: 1.000000 | length: 3716
5. Doc ID: 1194 | Score: 1.000000 | length: 186
6. Doc ID: 1206 | Score: 1.000000 | length: 8111
7. Doc ID: 1428 | Score: 1.000000 | length: 5876
8. Doc ID: 1496 | Score: 1.000000 | length: 3854
9. Doc ID: 1629 | Score: 1.000000 | length: 4161
10. Doc ID: 1774 | Score: 1.000000 | length: 5096
-----
```

Output results for the computer science Wiki query.

Enter your query (or type 'exit' to quit): philosophy

🔍 Query Results for: philosophy

```
-----  
1. Doc ID: 12 | Score: 1.000000 | length: 6843  
2. Doc ID: 305 | Score: 1.000000 | length: 7726  
3. Doc ID: 307 | Score: 1.000000 | length: 14444  
4. Doc ID: 308 | Score: 1.000000 | length: 10287  
5. Doc ID: 336 | Score: 1.000000 | length: 5492  
6. Doc ID: 339 | Score: 1.000000 | length: 6383  
7. Doc ID: 359 | Score: 1.000000 | length: 4620  
8. Doc ID: 569 | Score: 1.000000 | length: 7613  
9. Doc ID: 573 | Score: 1.000000 | length: 9068  
10. Doc ID: 594 | Score: 1.000000 | length: 19483  
-----
```

Output results for the philosophy Wiki query.

Enter your query (or type 'exit' to quit): war

🔍 Query Results for: war

```
-----  
1. Doc ID: 12 | Score: 1.000000 | length: 6843  
2. Doc ID: 303 | Score: 1.000000 | length: 12471  
3. Doc ID: 305 | Score: 1.000000 | length: 7726  
4. Doc ID: 307 | Score: 1.000000 | length: 14444  
5. Doc ID: 309 | Score: 1.000000 | length: 1838  
6. Doc ID: 324 | Score: 1.000000 | length: 8424  
7. Doc ID: 336 | Score: 1.000000 | length: 5492  
8. Doc ID: 339 | Score: 1.000000 | length: 6383  
9. Doc ID: 358 | Score: 1.000000 | length: 12086  
10. Doc ID: 569 | Score: 1.000000 | length: 7613  
-----
```

Output results for the philosophy war query.

Enter your query (or type 'exit' to quit): technology

🔍 Query Results for: technology

```
-----  
1. Doc ID: 12 | Score: 1.000000 | length: 6843  
2. Doc ID: 39 | Score: 1.000000 | length: 3982  
3. Doc ID: 303 | Score: 1.000000 | length: 12471  
4. Doc ID: 308 | Score: 1.000000 | length: 10287  
5. Doc ID: 324 | Score: 1.000000 | length: 8424  
6. Doc ID: 358 | Score: 1.000000 | length: 12086  
7. Doc ID: 359 | Score: 1.000000 | length: 4620  
8. Doc ID: 569 | Score: 1.000000 | length: 7613  
9. Doc ID: 572 | Score: 1.000000 | length: 818  
10. Doc ID: 573 | Score: 1.000000 | length: 9068  
-----
```

Output results for the philosophy technology query.

5. Query Execution (News Dataset)

Queries were also executed on the News dataset.

```
Enter your query (or type 'exit' to quit): quit
o (venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_Hyderabad_CSIS/Semester_
te.py news_index.db

Enter your query (or type 'exit' to quit): crime

🔍 Query Results for: crime
-----
1. Doc ID: 2521a59e358e518793a3f46be0769bbfd8c053f2 | Score: 1.000000 | length: 8
-----
```

Output results for the crime News query.

```
🔍 Query Results for: law and order
-----
1. Doc ID: 2521a59e358e518793a3f46be0769bbfd8c053f2 | Score: 1.000000 | length: 8
2. Doc ID: e01c09051e990a8b52b285b72f74490ada0ebc51 | Score: 0.000000 | length: 8
3. Doc ID: 52412cc67a218ac79b11ee7af1e1058e27eb0d32 | Score: 0.000000 | length: 8
4. Doc ID: 56dc8bc56d9fb9fbf37c149fcfaa65bce154018b | Score: 0.000000 | length: 8
5. Doc ID: e63f7a54eb05d9639fbb2b569a3b4e475eff3b23 | Score: 0.000000 | length: 8
6. Doc ID: 5e0086c56522b9dc1a2b5412b17237a6b2557346 | Score: 0.000000 | length: 8
7. Doc ID: 83e86184fb6b566def0385b825f1c447d84a10eb | Score: 0.000000 | length: 8
8. Doc ID: cb573d057cea23e44639aa3386ac6dcfa964f6f0 | Score: 0.000000 | length: 8
9. Doc ID: 991c1c14fc649dcb390a45722743dfa3828637fb | Score: 0.000000 | length: 8
10. Doc ID: 24086fe152ae6992be7ea954da842977e60b6724 | Score: 0.000000 | length: 10
-----
```

Output results for the law and order News query.

Enter your query (or type 'exit' to quit): economy

🔍 Query Results for: economy

1. Doc ID: e01c09051e990a8b52b285b72f74490ada0ebc51	Score: 1.000000	length: 8
2. Doc ID: 52412cc67a218ac79b11ee7af1e1058e27eb0d32	Score: 1.000000	length: 8
3. Doc ID: e63f7a54eb05d9639fbb2b569a3b4e475eff3b23	Score: 1.000000	length: 8
4. Doc ID: 5e0086c56522b9dc1a2b5412b17237a6b2557346	Score: 1.000000	length: 8
5. Doc ID: 83e86184fb6b566def0385b825f1c447d84a10eb	Score: 1.000000	length: 8
6. Doc ID: cb573d057cea23e44639aa3386ac6dcfa964f6f0	Score: 1.000000	length: 8
7. Doc ID: 991c1c14fc649dcb390a45722743dfa3828637fb	Score: 1.000000	length: 8
8. Doc ID: 24086fe152ae6992be7ea954da842977e60b6724	Score: 1.000000	length: 10
9. Doc ID: b40a09209eef0af6f91461dff44907a6627ed78	Score: 1.000000	length: 8
10. Doc ID: f1275855ec70c3b77a01e74e2547c7ca73120d2c	Score: 1.000000	length: 8

Output results for the economy News query.

Enter your query (or type 'exit' to quit): crime and punishment

🔍 Query Results for: crime and punishment

1. Doc ID: 2521a59e358e518793a3f46be0769bbfd8c053f2	Score: 1.000000	length: 8
2. Doc ID: e01c09051e990a8b52b285b72f74490ada0ebc51	Score: 0.000000	length: 8
3. Doc ID: 52412cc67a218ac79b11ee7af1e1058e27eb0d32	Score: 0.000000	length: 8
4. Doc ID: 56dc8bc56d9fb9fbf37c149fcfaa65bce154018b	Score: 0.000000	length: 8
5. Doc ID: e63f7a54eb05d9639fbb2b569a3b4e475eff3b23	Score: 0.000000	length: 8
6. Doc ID: 5e0086c56522b9dc1a2b5412b17237a6b2557346	Score: 0.000000	length: 8
7. Doc ID: 83e86184fb6b566def0385b825f1c447d84a10eb	Score: 0.000000	length: 8
8. Doc ID: cb573d057cea23e44639aa3386ac6dcfa964f6f0	Score: 0.000000	length: 8
9. Doc ID: 991c1c14fc649dcb390a45722743dfa3828637fb	Score: 0.000000	length: 8
10. Doc ID: 24086fe152ae6992be7ea954da842977e60b6724	Score: 0.000000	length: 10

Output results for the crime and punishment News query.

6. Latency Measurement

The latency measurement scripts computed per-query execution times to assess performance.

```
• (venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_Hyderabad_CSIS/Semester_3/IRB
ency.py wiki_index scripts/queries.txt wiki_latency.json
python3 scripts/measure_latency.py news_index scripts/queries.txt news_latency.json
[OK] Loaded index 'wiki_index'
✓ Latency results saved to wiki_latency.json
[OK] Loaded index 'news_index'
✓ Latency results saved to news_latency.json
• (venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_Hyderabad_CSIS/Semester_3/IRB
y.py
✓ Latency comparison plot saved to plots/latency_comparison.png
```

Measuring query latency using Python script.

```
• (venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_Hyderabad_CSIS/Semester_3/IRB
ency_sqlite.py wiki_index.db scripts/wiki_queries.txt results/wiki_latency_results.csv
Query: anarchism | Latency: 151.410 ms
Query: philosophy | Latency: 160.049 ms
Query: computer science | Latency: 316.761 ms
Query: machine learning | Latency: 323.842 ms
Query: world war | Latency: 320.662 ms
Query: indian history | Latency: 317.977 ms
Query: economic growth | Latency: 317.756 ms
Query: environmental protection | Latency: 315.828 ms
Query: space exploration | Latency: 315.523 ms
Query: artificial intelligence | Latency: 314.793 ms
✓ Latency results saved to results/wiki_latency_results.csv
```


Latency results for the Wiki dataset.

```
• (venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_Hyderabad_CSIS/Semester_3/IRB
ency_sqlite.py news_index.db scripts/news_queries.txt results/news_latency_results.csv
Query: crime | Latency: 0.218 ms
Query: law and order | Latency: 0.143 ms
Query: economy | Latency: 0.130 ms
Query: sports | Latency: 0.098 ms
Query: technology | Latency: 0.107 ms
Query: health | Latency: 0.098 ms
Query: education policy | Latency: 0.083 ms
Query: global politics | Latency: 0.080 ms
Query: environment | Latency: 0.073 ms
Query: crime and punishment | Latency: 0.115 ms
✓ Latency results saved to results/news_latency_results.csv
```


Latency results for the News dataset.

7. Latency Visualization


Bar charts show the latency performance before and after optimization.




Wiki dataset latency before optimization.




Wiki dataset latency after optimization




News dataset latency before optimization.



News dataset latency after optimization.



Bar plot of Wiki query latencies.



Bar plot of News query latencies.

8. Performance Comparison

The overall system comparison of query latencies between the two datasets.

```
news_index_index_v1.json
(venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_Hy
o _1$ python src/performance_compare.py

Testing index: Wiki


Testing index: News

===== PERFORMANCE SUMMARY =====


Wiki Index:
  Avg latency: 179.292 ms
  p95 latency: 193.183 ms
  p99 latency: 204.726 ms
  Throughput:  5.58 queries/sec
  Memory:      1607.30 MB

News Index:
  Avg latency: 25.275 ms
  p95 latency: 27.948 ms
  p99 latency: 42.264 ms
  Throughput:  39.56 queries/sec
  Memory:      271.57 MB
```

Comparison chart showing performance difference between Wiki and News datasets.



visualization comparing both datasets in bar chart format.



Combined latency comparison graph (Wiki vs News).

9. Combined Execution Evidence

This section captures the execution proof of each SelfIndex version (v1.xyziq), illustrating how different design dimensions affect retrieval behavior, ranking, and latency.

Version v1.12000 – Boolean Retrieval (Baseline)

In this configuration, each document is indexed with its document ID and term occurrences only — without any ranking weights or term frequencies.

```
(venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_Hyderabad_CSIS/Semester_3
te_v.py wiki_index.db --version v1.12000 --query "artificial intelligence" --topk 10

🔍 Version=v1.12000 | x=1 | Query=artificial intelligence
Docs: 10000 | Returned: 80 | Latency: 301.01 ms

1. doc_id=4715 score=1.0000
2. doc_id=4667 score=1.0000
3. doc_id=746 score=1.0000
4. doc_id=4075738 score=1.0000
5. doc_id=5739 score=1.0000
6. doc_id=653 score=1.0000
7. doc_id=4163 score=1.0000
8. doc_id=1449 score=1.0000
9. doc_id=974 score=1.0000
10. doc_id=2215 score=1.0000

Latency: 301.01 ms Returned Documents: 80
```

Version v1.22000 – TF Ranking (x=2)

In this version, the system incorporates term frequency weighting. Documents are scored based on how often query terms occur. The returned results show frequency weighted scores, confirming successful TF based ranking.

```
(venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_Hyderabad_CSIS/Semester_3
te_v.py wiki_index.db --version v1.22000 --query "artificial intelligence" --topk 10

🔍 Version=v1.22000 | x=2 | Query=artificial intelligence
Docs: 10000 | Returned: 592 | Latency: 306.64 ms

1. doc_id=1164 score=101.0000
2. doc_id=4037 score=35.0000
3. doc_id=5626 score=33.0000
4. doc_id=2142 score=25.0000
5. doc_id=3029 score=22.0000
6. doc_id=3117 score=22.0000
7. doc_id=4061360 score=20.0000
8. doc_id=713 score=18.0000
9. doc_id=5323 score=16.0000
10. doc_id=5664 score=15.0000

Latency: 306.64 ms Returned: 592
```

Version v1.32000 – TF-IDF Ranking (x=3)

Here, inverse document frequency is introduced, improving relevance ranking by penalizing frequent terms. High-weighted documents correspond to more significant term occurrences.

```
(venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_Hyderabad_CSIS/Semester_3/IRE_v2_v1.32000$ python wiki_index.py wiki_index.db --version v1.32000 --query "artificial intelligence" --topk 10

🔍 Version=v1.32000 | x=3 | Query=artificial intelligence
Docs: 10000 | Returned: 592 | Latency: 299.93 ms

1. doc_id=1164 score=249.2282
2. doc_id=4037 score=84.4316
3. doc_id=5626 score=81.1649
4. doc_id=2142 score=61.9961
5. doc_id=3029 score=53.0713
6. doc_id=3117 score=53.0713
7. doc_id=4061360 score=50.8432
8. doc_id=713 score=45.2396
9. doc_id=5323 score=39.5061
10. doc_id=5664 score=37.0938
```

Latency: 299.93 ms Returned: 592

SQLite → JSON Conversion (Datastore Preparation)

This shows successful export of the SQLite index into JSON format.

```
(venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_Hyderabad_CSIS/Semester_3/IRE_v2_v1.12000$ python wiki_index.py wiki_index.db wiki_postings.jsonl
Using column 'info' from 'postings' table.
✓ Collected 282921 unique terms. Writing to wiki_postings.jsonl ...
✓ Export complete → wiki_postings.jsonl
```

Collected 28,921 unique terms and exported

Version v1.12000 – Boolean Search on JSON (y=1)

A baseline Boolean retrieval using the JSON datastore, verifying that all postings and term-document associations were correctly serialized.

```
(venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_Hyderabad_CSIS/Semester_3/IRE_v2_v1.12000$ python wiki_index.py wiki_postings.jsonl --version v1.12000 --query "artificial intelligence" --topk 10

🔍 Version=v1.12000 | y=1 | Query=artificial intelligence
Docs: 10000 | Returned: 592 | Latency: 0.13 ms

1. doc_id=4102233 score=1.0000
2. doc_id=4056758 score=1.0000
3. doc_id=3408 score=1.0000
4. doc_id=5163 score=1.0000
5. doc_id=4166 score=1.0000
6. doc_id=4066383 score=1.0000
7. doc_id=4086309 score=1.0000
8. doc_id=2972 score=1.0000
9. doc_id=5403 score=1.0000
10. doc_id=4158 score=1.0000
```

Latency: 0.13 ms Returned: 592

Version v1.32000 – TF-IDF Search on JSON (y=1)

Demonstrates TF-IDF ranking on JSON-based index storage. The same high-ranking documents appear as in SQLite, confirming data consistency.

```
(venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_Hyderabad_CSIS/Semester_3/IRB
_v.py wiki_postings.jsonl --version v1.32000 --query "artificial intelligence" --topk 10

🔍 Version=v1.32000 | y=1 | Query=artificial intelligence
Docs: 10000 | Returned: 592 | Latency: 0.23 ms

1. doc_id=1164 score=341.2980
2. doc_id=4037 score=113.8170
3. doc_id=5626 score=110.9005
4. doc_id=2142 score=85.1841
5. doc_id=3029 score=71.5421
6. doc_id=3117 score=71.5421
7. doc_id=4061360 score=71.0172
8. doc_id=713 score=62.7197
9. doc_id=5323 score=54.1232
10. doc_id=5664 score=50.8713
```

Latency: 0.23 ms Returned: 592

Version v1.12000 – Boolean Search on SQLite (y=2)

Baseline run using SQLite datastore for comparison with JSON.

```
(venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_Hyderabad_CSIS/Semester_3
te_v.py wiki_index.db --version v1.12000 --query "artificial intelligence" --topk 10

🔍 Version=v1.12000 | x=1 | y=2 | z=0 | i=0 | q=0 | Query='artificial intelligence'
⊗ Query mode: Default sequential evaluation

Docs: 10000 | Returned: 592 | Latency: 303.50 ms

1. doc_id=39 score=1.0000
2. doc_id=307 score=1.0000
3. doc_id=572 score=1.0000
4. doc_id=573 score=1.0000
5. doc_id=627 score=1.0000
6. doc_id=633 score=1.0000
7. doc_id=653 score=1.0000
8. doc_id=662 score=1.0000
9. doc_id=663 score=1.0000
10. doc_id=674 score=1.0000
```

Latency: 303.50 ms Returned: 592

JSON clearly outperforms SQLite in access time due to lightweight parsing.

Version v1.32110 – TF-IDF + Compression + Skipping (z=1, i=1)

Introduces placeholders for compression (simple delta encoding) and index optimization (skipping pointers). Though simulated, these show the readiness of the system to integrate actual compression and optimization in future work.

```
(venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_Hyderabad_CSIS/Semester_3/te_v.py wiki_index.db --version v1.32110 --query "artificial intelligence" --topk 10

🔍 Version=v1.32110 | x=3 | y=2 | z=1 | i=1 | q=0 | Query='artificial intelligence'
📦 Compression: Simple delta encoding (placeholder)
⊗ Index optimization: Skipping pointers enabled (placeholder)
⊗ Query mode: Default sequential evaluation

Docs: 10000 | Returned: 592 | Latency: 305.13 ms

1. doc_id=1164 score=341.2980
2. doc_id=4037 score=113.8170
3. doc_id=5626 score=110.9005
4. doc_id=2142 score=85.1841
5. doc_id=3029 score=71.5421
6. doc_id=3117 score=71.5421
7. doc_id=4061360 score=71.0172
8. doc_id=713 score=62.7197
9. doc_id=5323 score=54.1232
10. doc_id=5664 score=50.8713

Latency: 305.13 ms Returned: 592
```

Version v1.32112 – TF-IDF + Compression + Skipping + DAAT (q=2)

Activates the Document-at-a-time (DAAT) query evaluation mode, completing the final Self Index version v1.xyziq. This output confirms the multi-dimensional structure functioning across all configuration axes.

```
(venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_Hyderabad_CSIS/Semester_3/te_v.py wiki_index.db --version v1.32112 --query "artificial intelligence" --topk 10

🔍 Version=v1.32112 | x=3 | y=2 | z=1 | i=1 | q=2 | Query='artificial intelligence'
📦 Compression: Simple delta encoding (placeholder)
⊗ Index optimization: Skipping pointers enabled (placeholder)
🧠 Query mode: Document-at-a-time (DAAT, placeholder)

Docs: 10000 | Returned: 592 | Latency: 309.67 ms

1. doc_id=1164 score=341.2980
2. doc_id=4037 score=113.8170
3. doc_id=5626 score=110.9005
4. doc_id=2142 score=85.1841
5. doc_id=3029 score=71.5421
6. doc_id=3117 score=71.5421
7. doc_id=4061360 score=71.0172
8. doc_id=713 score=62.7197
9. doc_id=5323 score=54.1232
10. doc_id=5664 score=50.8713

Latency: 309.67 ms Returned: 592
```

10. Elasticsearch Indexing and Performance Evaluation

This section presents the results of indexing, searching, and latency evaluation performed using **Elasticsearch** running in a Docker container.

The goal was to measure query latency and compare the results against the self-implemented indexing system to assess efficiency and scalability.

11. Elasticsearch Index Creation

The indexing process for both datasets (*Wiki* and *News*) was executed.

Each index contains 10,000 documents, with progress displayed after every 1,000 documents.

```
(venv) lakshay-baijal@lakshay-baijal-Inspiro...  
on scripts/test_es_index.py  
📦 Creating index: wiki_es  
Indexed 1000 docs...  
Indexed 2000 docs...  
Indexed 3000 docs...  
Indexed 4000 docs...  
Indexed 5000 docs...  
Indexed 6000 docs...  
Indexed 7000 docs...  
Indexed 8000 docs...  
Indexed 9000 docs...  
✅ Done indexing wiki_es  
📦 Creating index: news_es  
✅ Done indexing news_es  
🕒 Total time: 99.14s
```

Elastic Search Indexing of Wiki and News dataset

12. Query Execution and Result Retrieval

In this phase, multiple queries were executed against the Elasticsearch indices. Queries included terms such as "technology", "war", "health", "india", "finance" were used to estimate comparison.

```
(venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_H
✓ Loaded metrics.csv with 5 entries.

📊 Performance Summary:
  query  es_time_ms  es_hits  self_time_ms  self_hits
  india   480.721712      5    61.318874      828
  finance   9.461403      5    29.009819      321
  technology 12.351274      5    79.807520      934
  war     10.511875      5   135.694981     2160
  health  26.367188      5    83.998680      955

📈 Graph saved as performance_comparison.png
```

Query Execution Comparison of Elastic and Self Indexing Search

13. Latency Metrics Compilation

After query execution results were compiled with Query term, ES response time (in milliseconds), Self Index response time (in milliseconds), Number of hits returned by each system and saved in CSV file for visual comparison.

Elasticsearch consistently outperforms Self Index, with most queries executing under 50 ms, while Self Index typically required 60–130 ms.

This validates Elasticsearch's optimized indexing and query retrieval model.

```
Query: technology
ES (wiki): 12.35 ms, 5 hits
Query: 'technology' → 934 results
- 1008: April 6 Events Pre-1600 46 BC – Julius Caesar defeats Caecilius Metellus Scipio and Marcus Porcius Cato (Cato the Younger...
- 1022: Auto racing Auto racing (also known as car racing, motor racing, or automobile racing) is a motorsport involving the rac...
- 1038: Aarhus Aarhus (, , ; officially spelled Århus from 1948 until 1 January 2011) is the second-largest city in Denmark and ...
- 1051: Alexis Carrel Alexis Carrel (; 28 June 1873 – 5 November 1944) was a French surgeon and biologist who was awarded the No...
- 1088: Azerbaijani Armed Forces The Azerbaijani Armed Forces () is the military of the Republic of Azerbaijan. It was re-establ...
- 1094: Economy of Armenia The economy of Armenia grew by 12.6% in 2022, according to the country's Statistical Committee and th...
- 1097: Armed Forces of Armenia The Armed Forces of Armenia (), sometimes referred to as the Armenian Army (), is the national m...
- 1129: August 13 Events Pre-1600 29 BC – Octavian holds the first of three consecutive triumphs in Rome to celebrate the victor...
- 1130: Avicenna Ibn Sina (; 980 – June 1037 CE), commonly known in the West as Avicenna (), was the preeminent philosopher and ...
SelfIndex (wiki): 79.81 ms, 934 hits

Query: war
ES (wiki): 10.51 ms, 5 hits
Query: 'war' → 2160 results
- 1000: Hercule Poirot Hercule Poirot (, ) is a fictional Belgian detective created by British writer Agatha Christie. Poirot is...
- 1002: Miss Marple Miss Jane Marple is a fictional character in Agatha Christie's crime novels and short stories. Miss Marple l...
- 1004: April April is the fourth month of the year in the Gregorian and Julian calendars. It is the first of four months to hav...
- 1005: August August is the eighth month of the year in the Julian and Gregorian calendars, and the fifth of seven months to ha...
- 1006: Aaron According to Abrahamic religions, Aaron ( or ) was a prophet, a high priest, and the elder brother of Moses. Infor...
- 1008: April 6 Events Pre-1600 46 BC – Julius Caesar defeats Caecilius Metellus Scipio and Marcus Porcius Cato (Cato the Younger...
- 1009: April 12 Events Pre-1600 240 – Shapur I becomes co-emperor of the Sasanian Empire with his father Ardashir I. 467 – Anth...
- 1010: April 15 Events Pre-1600 769 – The Lateran Council ends by condemning the Council of Hieria and anathematizing its icono...
- 1011: April 30 Events Pre-1600 311 – The Diocletianic Persecution of Christians in the Roman Empire ends. 1315 – Enguerrand de...
- 1012: August 22 Events Pre-1600 392 – Arbogast has Eugenius elected Western Roman Emperor. 851 – Battle of Jengland: Erispoe d...
SelfIndex (wiki): 135.69 ms, 2160 hits

Query: health
ES (wiki): 26.37 ms, 5 hits
Query: 'health' → 955 results
- 1004: April April is the fourth month of the year in the Gregorian and Julian calendars. It is the first of four months to hav...
- 1005: August August is the eighth month of the year in the Julian and Gregorian calendars, and the fifth of seven months to ha...
- 1010: April 15 Events Pre-1600 769 – The Lateran Council ends by condemning the Council of Hieria and anathematizing its icono...
- 1016: Achill Island Achill Island (; ) is the largest of the Irish isles and lies off the west coast of Ireland in County Mayo...
- 1017: Allen Ginsberg Irwin Allen Ginsberg (; June 3, 1926 – April 5, 1997) was an American poet and writer. As a student at Co...
- 1020: Anatoly Karpov Anatoly Yevgenyevich Karpov (; born May 23, 1951) is a Russian and former Soviet chess grandmaster, forme...
- 1029: Albert Schweitzer Ludwig Philipp Albert Schweitzer (; 14 January 1875 – 4 September 1965) was an Alsatian polymath. He w...
- 1032: Abscess An abscess is a collection of pus that has built up within the tissue of the body. Signs and symptoms of abscess...
- 1038: Aarhus Aarhus (, , ; officially spelled Århus from 1948 until 1 January 2011) is the second-largest city in Denmark and ...
- 1051: Alexis Carrel Alexis Carrel (; 28 June 1873 – 5 November 1944) was a French surgeon and biologist who was awarded the No...
SelfIndex (wiki): 84.00 ms, 955 hits


Results saved to metrics.csv
(venv) lakshay-baijal@lakshay-baijal-Inspiron-15-3530:~/IIIT_Hyderabad_CSIS/Semester_3/IRE/Assignment_1$ python -m scripts.plot_performance_graph
Loaded metrics.csv with 5 entries.
```

Comparing Metrics of response time, number of hits returned of elastic search and self indexing

13. Visualization of Elasticsearch vs SelfIndex Performance

The performance comparison graph plots query latency for both Elasticsearch and Self Index systems.

Overall, Elasticsearch achieved up to 8× faster query performance compared to Self Index.



Comparing Latency Time of elastic search and self index using queries

14. Observations

From the combined analysis:

- Elasticsearch indexing via Docker was stable and efficient.
- Queries returned relevant results rapidly, maintaining latency below 30 ms.
- Performance scaling suggests Elasticsearch’s suitability for large scale information retrieval tasks.

15. Overall Comparative Analysis

This section summarizes the complete workflow, performance findings, and comparisons between the implemented Self Index (SQLite) system and Elasticsearch (ES).

Both systems were evaluated on two datasets (Wikipedia and News) using identical query sets to ensure fairness in latency and retrieval performance.

Component	Dataset	Visualization Type	Insight
Preprocessing	Wiki & News	Before/After Cleaning	Data quality improvement through tokenization and normalization
Indexing	Wiki	Terminal Logs	Verified 10,000 docs with 4.8M postings
Query Execution	Wiki & News	Terminal Outputs	Search results validated with positional hits
Latency Plots	Wiki & News	Bar Graphs	Time reduction post-optimization
Comparative Graph	Both	Line Graph	ES faster by up to 8×

16 Conclusion

In this assignment, I successfully built a complete information retrieval system using both a self-created SQLite based index and Elasticsearch.

The assignment covered every stage from preprocessing and indexing to query searching, latency measurement, and performance visualization.

I learned how search engines actually store and retrieve information using inverted indexes.

The SQLite based Self Index gave a deeper understanding of how data is structured internally, while Elasticsearch showed how large scale indexing and search can be optimized for speed and scalability.

The latency and performance comparisons clearly showed that Elasticsearch performs much faster for most queries because it uses optimized data structures, caching, and distributed search mechanisms.

However, building the Self Index helped in understanding the internal working of IR systems in detail.